

Research Statement

Vineet Kumar, *Yale School of Management*

My overall research focus is at the intersection of technology and business & society.¹ Digital technology has had substantial impact on consumers, firms and society. Given my background in both engineering and business, I am especially interested in two focus areas.

- (A) *Digital business models* have specific features relative to the markets for commonly studied goods. My research focuses on strategic long-run choices relevant to digital firms, a topic of interest to multiple fields. I study this area using methods based on microfoundations of agents' preferences, typically empirical structural models.
- (B) Building *theory-based machine learning (ML)* methods incorporating structured knowledge (theory) developed from first principles to be human interpretable, and provide representations satisfying required properties (e.g. monotonicity). These methods also enable new business models (creating new sources of value).

(A) Digital Business Models

The business model of a firm denotes a set of integrated strategic choices that drive performance in the marketplace. With traditional products, these models are well established; digital offers new possibilities (e.g. freemium, versioning, social etc.) that deserve careful study; these choices and their alignment provide the foundation for success or failure.

Within digital business models, my research can broadly be themed as focusing on three different sources of connections: between *products*, *consumers* and *data*.

1) *Connections across Products – Product Line*

I detail the connections between products, with a research overview and agenda provided in [P1].

Bundling Hardware and Software: In [P4], I examine the dynamic effects of bundling in digital platforms (hardware+software). Should a firm create mixed bundles (products+bundles), pure bundles, or no bundles? I investigate the dynamics of bundling using data on sales and product characteristics in markets for videogame consoles (hardware) and game titles (software). In dynamic settings with intertemporal tradeoffs, bundling is more effective with *positive correlation* of valuation across products. Bundles act by pulling demand forward, i.e. consumers buy now rather than wait, a new mechanism, and higher sales obtain with positive correlation. Prior to this research, bundling was studied in a static setup and found to be more effective with negative correlation. I also show that bundling is more effective when indirect network effects are weaker. I develop a novel identification strategy for correlation in valuations based on the tying ratio, leveraging the market feature that consumers purchase one hardware unit, but many software units.

Versioning: Freemium is the most popular digital business model (e.g. app stores, cloud storage,

¹For simplicity of exposition, this document is written stylistically in the first person singular, although most of my work is in collaboration with a wonderful set of co-authors.

SaaS). I provide an overview of issues in Freemium for a general audience in [P11]. Few consumers (typically <5%) migrate to the paid version, so product design is of critical importance. I undertake a deep dive into designing freemium for a storage service, examining the growth-monetization tradeoff. In [P16], I evaluate product design (value of free version) and referral incentives, which impact the value of the free product, and therefore, upgrades. In counterfactuals, I find that offering a greater referral bonus relative to the firm's bonus can increase growth but reduce monetization, while beyond a threshold, it can decrease both growth and monetization. I study how to optimally structure dynamic referral incentives.

I examine the strategic use of time as a versioning and monetization strategy (wait for free) in [P18]. This strategy is commonly employed by platforms publishing content (e.g. novels as serialized episodes), with consumers having complementary value for content across episodes. I leverage a natural experiment – wait time is reduced by the platform – to study the impact on downstream consumer choices. Existing consumers were found to increase paid consumption, and more new consumers start reading, increasing aggregate consumption. I show how a microfounded mechanism with complementarity can rationalize these data patterns. Overall, this demonstrates the strategic role of temporal versioning in driving choices, which has not been shown earlier.

Digital Transformation: In [P10], I study digital transformation strategy for a firm moving from one offering to another, typically a product based on an older technology (physical) to one based on newer technology (digital). I examine pricing and product design in this empirical setting. Using a panel from the “Netflix” model, I obtain consumer preferences for viewing content in physical form, and evaluate optimally pricing the product line. The product line contains multiple versions, which allow the firm to extract surplus from heterogeneous consumers. However, as the service time is reduced, the differentiation across versions also decreases. In counterfactuals, I uncover novel mechanisms demonstrating how improved service time (better for all customers) could result in lower profits and even lower revenue for the firm under optimal pricing, a previously unappreciated transformation risk.

Open Source: I examine the market for open source software [P14], where products made by competing firms share common elements. Open source contributions made either by developers or by any firm are available to all competing firms (e.g. Linux or Android). It is puzzling that encouraging free-riding can lead to high quality products; my model with interconnected markets (developer and product) explains how this happens. Developers signal their capabilities by making contributions of features to the open source software (public good). Free-riding is sustained in equilibrium since firms can build on features to differentiate on another complementary dimension (usability); I find that appropriability can even increase product quality. Broadly, this research provides insight into how these open source contributions impact the product market.

2) Connections across Data – Linking Purchase and Usage

The third connection that I examine is the linkage between types of data, i.e. purchase data and usage data. In digital settings, usage data is uniquely available, and valuable in obtaining insights about consumer preferences. However, most studies in marketing and economics involve only

purchase data. I connect usage and purchase data across several settings, including [P16, P10]. In [P3], I demonstrate how usage data is conceptually distinct, and is critically important for identification, not just estimation. Specifically, I show that identification of the Willingness to Pay (WTP) distribution for subscriptions without price variation is possible by leveraging usage data, but impossible without it. The key insight is that combining high-frequency usage data with purchase data allows for a conceptual leap in identification of the distribution of consumer willingness-to-pay (WTP), which was not thought to be possible without price variation. I combine usage data with exogenous factors impacting usage to first estimate usage utility, then aggregate this stream, combined with purchase data to obtain the WTP distribution. I can then conduct counterfactual analyses, such as product design. The framework is flexible in accommodating a large class of usage utility models, making it widely applicable.

3) Connections across Consumers – Networks

While word-of-mouth effects impact business models, this research is broadly applicable to a wide class of network interventions.

I investigate privacy-sensitive methods for leveraging network structure to obtain higher-degree nodes in unknown networks [P15], e.g. for word-of-mouth. The literature focuses on obtaining the entire network structure (not privacy-friendly!), and also does not offer provable guarantees. I examine two strategies (ego-based and alter-based), based on the friendship paradox, asking individuals to nominate random friends. I show that these strategies have distinct mathematical properties, and also propose a new network property called Inversity, which determines which strategy obtains better connected seeds, based on network structure. The strategies are simple to implement and offer provable guarantees of obtaining higher-degree individuals.

I empirically study whether using friendship paradox strategies can achieve greater adoption in [P13]. Using a model of communication and adoption across social networks estimated on data, I evaluate counterfactual seeding strategies. I show that ego-based friendship seeding outperforms random seeding, and surprisingly, obtains higher adoption than even leader-based seeding. The results hold across a range of specifications and networks, demonstrating robustness and empirical value, and have implications for referral design. In contrast to prior work, this research proves that it is possible to leverage networks to impact interventions with guarantees for *any* network, without knowing the network.

Methodological Overview and Contributions: Theory or structured knowledge is central to my microfounded models. This structural approach yields estimates with clear economic interpretations, enabling the counterfactual evaluation of firm or regulatory policies. During the work on some of the above projects, I investigated methods for estimating dynamic demand models in technology markets. Examining the commonly used inclusive value approach, I show that it can lead to highly biased estimates of economically important quantities like elasticity and profits [P5]. Motivated by the need for models to flexibly accommodate dynamics with large state spaces, I develop a new method with lower computational complexity that can be easily estimated [P2].

Specifically, for a large class of problems (with terminating or renewal choice) using market-level data, it obtains consumer preferences with the computational complexity of a linear regression. Identification is proven formally and the method can be used for counterfactual analysis.

(B) Theory-based Machine Learning

My research focus in ML is based on integrating structured knowledge to develop better ML methods, which in turn, enables us to obtain insights about consumer responses. First, there is a growing set of important research questions where ML is required. Marketing in practice involves important elements that impact consumers that cannot be appropriately captured by structured data (e.g. text, visual design, music, videos). Yet, the vast majority of research has focused on structured data, since they are more available and have standardized methods to analyze them. Traditional non-ML quantitative methods are not great at capturing the nuances of unstructured data. Moreover, ML also makes it feasible to generate novel unstructured data like text or images.

Given their growing capabilities, ML models are being increasingly used in academia and industry. However, they are typically opaque black box models (e.g. ChatGPT), leading to significant problems. First, these models are highly complex (with billions of parameters) and humans don't know the interpretation of these parameters. We also don't know the data on which they are trained. Second, they do not have a true understanding of the consumer.² Third, they are typically not interpretable – we don't know *why* they work. They are prone to failure (e.g. hallucinations), and we cannot know *when* this will happen because we have little visibility.³

My focus is on solving these challenges, developing better ML methods based on theory or structured knowledge.⁴ My background, being fluent with both microfounded theory-based models and in ML, has enabled me to bring a unique perspective to this research. My research is focused around three aspects: (a) developing methods to incorporate structured knowledge into ML models, (b) provide complete model and data transparency, and (c) improve explainability and interpretability along with performance. Rather than using commercially available black box models, I develop models from basic elements. All of my ML research is transparent with the open source code publicly available for others to examine, critique and build on. I believe this approach can improve stakeholder trust and acceptance of AI systems.

The sources of knowledge that I examine include ideas and concepts from fields like marketing and economics, but is not restricted to these. Such knowledge can be used across both structured and unstructured data. However, building in theoretical foundations within ML methods is typically quite challenging. The question is how we can incorporate theory into complex deep learning models with unstructured data. Such a challenge holds even in the case of reinforcement learning with structured data. Below, I detail the specifics of my research in ML to illustrate this approach.

Visual Characteristics: I aim to quantify consumer preferences for visual characteristics in [P19].

²The same CNN deep net models used for marketing applications are also used in biology.

³Consider the algorithms for self-driving vehicles illustrating these points – only now are we beginning to get some idea about the inner workings. See [The Hidden Autopilot Data That Reveals Why Teslas Crash](#) (Wall Street Journal, 07/30/2024).

⁴Other researchers are trying to integrate theory from their academic fields into ML, e.g. in physics [O1].

Visual appearance is high-dimensional and hard to characterize and explain, without human input. I develop a theory-based algorithm to automatically discover and quantify visual characteristics of products. Theory plays a crucial role in the following ways. First, the objective function is designed to incentivize low-dimensional and orthogonal representations, based on the idea that the product designs satisfy that property. Second is the idea that products often have a distinct look tied to the brand, with recognizable “visual signatures” like LV handbags or BMW cars. Brands typically have a consistent aesthetic, and consumers form expectations around this. The method extracts visual characteristics from product images using brand and other characteristics to supervise disentanglement. The method obtains significantly better performance *and* interpretability, and importantly, can *generate* counterfactual visual designs without human judgment.

Music and Emotion: In the music emotion research [P7], the ideas about consonance and dissonance of music, and how that connects to the listener’s emotion is aided by using domain knowledge (theory) from multiple fields. Specifically, the knowledge is based on both: (a) the mathematics of sound waves and (b) psychology of human music perception. Explaining why a listener feels a specific emotion when listening to music was a black box earlier. Here, theory is used as the basis for creating flexible and non-contiguous consonance filters, helping obtain a representation that enables explainability, so we can visualize the features of music impacting listener responses. The application also enables a form of contextual targeting for ads, without using any consumer data (aiding privacy).

Learning Unknown Demand Curves: To learn an unknown demand curve by experimentation, I develop a reinforcement learning model with nonparametric multi-armed bandits (MAB) [P20]. Economic theory informs us that demand curves are downward sloping.⁵ The classic experiment (A/B testing) benchmark is inefficient since it explores all prices equally, whereas MABs provide a more sophisticated approach (“learning while earning”). However, *how* to incorporate theory here is not quite obvious. I find that monotonicity adds two sources of value. First, it improves the performance of the algorithm substantially since it learns not just from each individual price (or arm) experimented, but across arms (an informational externality). Second, including theoretical guarantees that the resulting demand curve obtained is monotonic is especially important when algorithms are trusted and used to make automated pricing decisions. Without monotonicity, we can get an upward sloping demand curve, and the algorithm would choose unrealistically high prices, increasing the risk of failure.

Overall, all the ML-based research I have undertaken brings in the power of structured knowledge to enhance the capability of ML along several aspects: performance, interpretability / explainability, and providing representations satisfying desirable properties. In turn, these ML methods are useful to help us gain valuable insights into consumer and firm behavior, leading to a symbiotic process. I plan to continue to explore theory-based ML methods and new business models enabled by ML.

⁵There are exceptions (e.g. Veblen goods).

Teaching

I have developed and taught elective courses at the masters and doctoral levels, and have also contributed to the executive MBA and executive education. I use a mix of lectures to help with learning principles, complemented with discussions featuring case studies by the Socratic method. In *Digital Strategy*, I explore digital business models (e.g. Dropbox), and evaluate models used by disruptors and complementors. I investigate the drivers and barriers of platform success, and examine digital transformation and emerging technologies. Given the connection to my research, I've incorporated exercises based on research into the course material. Digital transformation has been used in custom executive education programs at Yale, e.g. YGELP.

I recently developed and taught a masters-level elective course titled *Artificial Intelligence: Strategy & Marketing*. The objective is to help students understand the ideas, gain familiarity with the methods and their business applications. I introduce the primary ML methods (supervised, supervised, reinforcement and generative), followed by an examination of how organizations obtain value from AI, through case studies. I also demonstrate how fairness (and bias) can be and are often designed into algorithms, whether intentionally or not, and the resulting implications for all stakeholders. In my pedagogical experiments, I've found that assigning students to present (carefully selected) academic papers can be effective, with faculty guidance.

References

- [P1] Steve Berry, Ahmed Khwaja, Vineet Kumar, Andres Musalem, Kenneth C Wilbur, Greg Allenby, Bharat Anand, Pradeep Chintagunta, W Michael Hanemann, Przemek Jeziorski, et al. Structural models of complementary choices. *Marketing Letters*, 25:245–256, 2014.
- [P2] Cheng Chou, Tim Derdenger, and Vineet Kumar. Linear estimation of aggregate dynamic discrete demand for durable goods: Overcoming the curse of dimensionality. *Marketing Science*, 38(5):888–909, 2019.
- [P3] (*) Cheng Chou and Vineet Kumar. Estimating Demand for Subscription Products: Identification of Willingness to Pay Without Price Variation. *Marketing Science*, 2024.
- [P4] (*) Timothy Derdenger and Vineet Kumar. The Dynamic Effects of Bundling as a Product Strategy. *Marketing Science*, 32(6):827–859, 2013.
- [P5] Timothy Derdenger and Vineet Kumar. Estimating dynamic discrete choice models with aggregate data: Properties of the inclusive value approximation. *Quantitative Marketing and Economics*, 17(4):359–384, 2019.
- [P6] Hortense Fong, Vineet Kumar, Anay Mehrotra, and Nisheeth K Vishnoi. Fairness for AUC via Feature Augmentation. *arXiv preprint arXiv:2111.12823*, 2021.
- [P7] (*) Hortense Fong, Vineet Kumar, and K Sudhir. A Theory-based Interpretable Deep Learning Architecture for Music Emotion. *Marketing Science (Forthcoming)*, 2024.
- [P8] Soheil Ghili, Vineet Kumar, and Fei Teng. Spatial Distribution of Access to Service: Theory and Evidence from Ridesharing. *Available at SSRN 4915262*, 2023.
- [P9] Pavel Kireyev, Vineet Kumar, and Elie Ofek. Match your own price? self-matching as a retailer’s multichannel pricing strategy. *Marketing Science*, 36(6):908–930, 2017.
- [P10] (*) Vineet Kumar and Yacheng Sun. Designing pricing strategy for operational and technological transformation. *Management Science*, 66(6):2706–2734, 2020.
- [P11] Vineet Kumar. Making ‘freemium’ work. *Harvard Business Review*, 92(5):27–29, 2014.
- [P12] Vineet Kumar and Kannan Srinivasan. Predicting Customer Value using Clumpiness – Commentary. *Marketing Science*, Mar-Apr, 2015.
- [P13] Vineet Kumar and K Sudhir. Can Friends Seed More Buzz? *Management Science (Accepted)*, 2024.
- [P14] Vineet Kumar, Brett R Gordon, and Kannan Srinivasan. Competitive Strategy for Open Source Software. *Marketing Science*, 30(6):1066–1078, 2011.
- [P15] Vineet Kumar, David Krackhardt, and Scott Feld. On the Friendship Paradox and Inversity: A network property with applications to privacy-sensitive network interventions. *Proceedings of the National Academy of Sciences*, 121(30), 2024.
- [P16] Clarence Lee, Vineet Kumar, and Sunil Gupta. Designing Freemium: Strategic Balancing of Growth and Monetization. *Available at SSRN 2767135*, 2019.
- [P17] Peter S Lee, Vineet Kumar, and K Sudhir. Intertemporal Price Discrimination for Serialized Media Products. *Working Paper*, 2024.
- [P18] Peter S Lee, Vineet Kumar, and K Sudhir. Monetizing Serialized Content: How Wait for Free Impacts Paid and Free Consumption. *Working Paper*, 2024.
- [P19] (*) Ankit Sisodia, Alex Burnap, and Vineet Kumar. Generative Interpretable Visual Design: Using Disentanglement for Visual Conjoint Analysis. *Journal of Marketing Research (Forthcoming)*, 2024.
- [P20] Ian Weaver, Vineet Kumar, and Lalit Jain. Nonparametric Bandits Leveraging Informational Externalities to Learn the Demand Curve. *Available at SSRN 4151019*, 2024.

Note: () papers chosen for distribution.*

Other References

- [O1] George Em Karniadakis et al. Physics-informed machine learning. *Nature Reviews Physics*. 3(6). 2021