



Course Information

Course Number:	MGT 853
Course Title:	Artificial Intelligence: Strategy + Marketing
Term and Year:	Spring-2 2024
Class Meeting Time, Day:	Section 01: Tue / Thu 2:40 pm – 4:00 pm, Evans Hall 4210 Section 02: Tue / Thu 4:10 pm – 5:30 pm, Evans Hall 4210

Contact Information

Professor:	TA(s):
Vineet Kumar	Seung Yoon Lee (01/02)
Office Location: 5455 Evans Hall	seungyoon.lee@yale.edu
Telephone Number: 203.436.9657	
E-mail Address: vineet.kumar@yale.edu	Peter Lee (01/02)
Office Hours: By Appointment (e-mail me)	peter.s.lee@yale.edu

Course Materials

- Textbook(s): None required
- Recommended Books: Instructor will provide required book chapters via Canvas.
- Readings: **See Outline of Class Sessions.**
- Software: **R or Python (or another language) may be used for assignments and project if needed. We will provide code for you to explore and comment on and potentially modify.**

Course Description and Objectives

Disclaimer: Syllabus is tentative. Please check latest version on Canvas by “Last Updated” date.

Course Description:

Artificial Intelligence is a general-purpose technology which has the potential to transform many aspects of business and society. In business, the impact ranges from commonplace predictive improvements at one end of the spectrum to opportunities for creating entirely new markets at the other. As background, the course will briefly introduce students to Artificial Intelligence / Machine Learning methods comprising of Unsupervised, Supervised and Reinforcement Learning. Through a combination of lectures and case studies, we will learn how to integrate AI into decision making, focusing on the strategic choices firms face in developing and using AI / ML, including generative AI. We will seek to understand how consumers, decision-makers other stakeholders evaluate decisions

made by AI systems and are impacted by it. We will examine the state-of-the-art in explainable and interpretable AI, which aim to make black-box decisions more transparent. Finally, the course will also explore issues at the intersection of AI and Society including fairness and bias, which have proved especially challenging to address.

Logistics: The course will be taught *in-person only* (there is no virtual option as per SOM policy), unless public health concerns require YSOM to change course.

Audit Policy: This is an elective class designed for graduate students at the Yale School of Management. Please note that this course will not be available for audit, and students must be officially registered to attend.

Course Objectives:

The course is designed to provide a strategic perspective on AI technologies, but given the broad scope of the topic, I've had to make tradeoffs in selecting areas to examine in depth. My goal is that a student who completes the course successfully should be able to:

- Understand the basics of AI and ML models
- Determine how AI objectives connect to business objectives and strategy
- Understand a framework for decisions on AI / ML and identify the major resources required to implement the chosen AI strategy
- Develop a perspective regarding new emerging AI technologies and how they could reshape markets and firms
- Evaluate the broader societal implications of AI, and how different stakeholders (consumers, employees, firms, regulators, investors and others) are impacted by AI.

Note: This is NOT a computer science or ML course. We will briefly review selected ML algorithms; however, we will not have a comprehensive or in-depth examination of how these algorithms work.

Course Requirements

Course Component	Details	Points
Assignments (Individual and Pairs)	See details below	50
Attendance and Class Participation (Individual)	Every class	30
Paper Presentation (Group)	See details below	25

Please see the Yale SOM Grading Policy at <http://portal.som.yale.edu/page/grading-policy>

Class Participation: AI Strategy is a class where significant value is derived from in-class discussions. You will be encouraged, and likely cold-called to air your views in class. Evaluation will be on the quality and to a lesser degree, the number of comments you make. Please come prepared to each class by reading the assigned material and thinking carefully about the preparation questions.

Description of Assignments and Projects

Assignments: Assignments are either individual, pair or group assignments. Please indicate the contribution of each person for assignments. Everyone in the group is expected to make a significant contribution to **each** assignment, so please budget time to discuss assignments. Late assignments will not be graded. No extensions or exceptions to any due dates should be expected. We don't have any makeup assignments for missed classes or assignments. However, if you have an extenuating circumstance, please contact **AASL**, and they will determine an appropriate course of action.

Your assignment submissions must be 4 pages or fewer (excluding any exhibits), in 11 or 12-point text font and double-spaced text. Make sure to include everyone's names in the submitted file. Please convert everything to one PDF file before submission and submit through Canvas. Please name the file in the following format. For individual assignments, please use the notation A1_LastName_FirstName.pdf. For pairs, include both names and only one submission. Details of Assignments will be posted on Canvas.

Due dates are listed in "Outline of Class Sessions" below. If you have any questions at all, please don't hesitate to check with the professor or TAs.

Project Presentation assignment (Group): All groups should be prepared to present in class. It is important to submit Slides through Canvas, so we can load them up well before class. Groups *must* be formed within section, not across sections.

After forming a group, students select a research paper to present at the end of the course. We will provide a curated list of academic papers representing cutting edge and seminal research in ML / AI to select from. In all cases, groups will receive the same presentation duration in class.

Custom Project: Students may also check with the instructor if they want to do a custom project. *Note:* The same or similar project(s) cannot be submitted or presented in multiple courses. Please check with the professor if you are not sure.

Meeting with Professor: You will need to set up a meeting with the professor / TAs to go over the paper and obtain feedback and approval. We will send out links to help schedule meeting times during predefined blocks.

The deliverable for the final project will be a slide deck designed to be submitted and presented at the end of the course. Students will be evaluated on both the presentation and the content in the slides. The instructor will discuss project details in class.

Projects will be evaluated on the following aspects: (a) clarity of thought and presentation, (b) depth of analysis, (c) insights obtained and (d) value to audience. Depth is valued more than a comprehensive but superficial presentation. Interpreting the model, results and learnings and providing your own perspective is valued more, rather than merely repeating what the paper says.

Yale SOM Policies

Please see the Yale School of Management Bulletin at <https://bulletin.yale.edu/bulletins/som/rights-and-responsibilities-students> for Rights and Responsibilities of students and for the honor code.

Policy on use of Large Language Models (LLMs)

In this course, you may use LLMs, provided you include in your submission the LLM used and the specific prompts that were asked. I view it as a useful way to learn things. However, you need to answer the assignments in your own words.

Laptop/Device Policy

Usage NOT allowed without the express permission of the instructor. There will be (at most) a couple of sessions where laptops will be useful in class, and the instructor will specifically mention these in class.

Outline of Class Sessions

#	Date	Topic	Assignment Due (9 am)
Module A: AI Foundations			
1	Mar 26 (Tue)	Course Introduction and Supervised and Un-supervised Algorithms	
2	Mar 28 (Thu)	ML Essentials	
3	April 2 (Tue)	Deep Learning, Reinforcement Learning and Generative Models	A1 (Individual / Pairs)
Module B: AI Decision Making Framework			
4	April 4 (Thu)	Economics of AI \iff Business Strategy	
5	April 9 (Tue)	Decision Making with AI / Interpretable and Explainable AI	A2 (Individual)
6	April 11 (Thu)	Ethical Issues in AI	Group Project Overview (one paragraph)
Module C: AI in Business + Society			
7	April 16 (Tue)	Uber (CASE)	A3 (Individual)
8	April 18 (Thu)	Zebra Medical (CASE)	
9	April 23 (Tue)	Miroglio Fashion (CASE)	A4 (Individual)
10	April 25 (Thu)	Human Capital with GROW	
11	April 30 (Tue)	Guest Speaker (awaiting confirmation)	
Module D: Project Presentations and Course Wrap			
12	May 2 (Thu)	Presentations	Presentation Slides Due for ALL groups on May 2
13	May 7 (Tue)	Presentations and Course Wrap	

Note: The content of some sessions is likely to change. I will post an announcement when there are material updates to the syllabus.

Detailed Outline of Class Sessions

Please note that “Preparation Questions” below are provided to help you prepare better for class discussion. These are NOT assignment questions and do not need to be turned in. Assignment questions will be posted separately on Canvas.

Readings available only on Canvas are marked with a (*).

Session 1: Introduction

Learning Material:

- (a) (*) The Business of AI, by E. Brynjolfsson and A. McAfee, HBR
- (b) (*) Artificial Intelligence for the Real World, T. Davenport; R. Ronanki, HBR
- (c) (*) Competing in the age of AI, by M. Iansiti and K. R. Lakhani, HBR
- (d) An Executives Guide to AI, by M. Chui, V. Kamalnath, and B. McCarthy, McKinsey & Co.
<https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/an-executives-guide-to-ai>
- (e) (Optional) The Building Blocks of an AI Strategy, by A. Joshi and M. Wade, MIT Sloan Review
<https://sloanreview.mit.edu/article/the-building-blocks-of-an-ai-strategy/>
- (f) (Optional) The Building Blocks of Artificial Intelligence, by P. Gerbert, M. Hecker, S. Steinhäuser, and P. Ruwolt
<https://www.bcg.com/publications/2017/technology-digital-strategy-building-blocks-artificial-intelligence>

Preparation Questions:

- (1) Where do you think the most promising applications of AI are in business? What criteria did you use in determining this?
- (2) Where do you think AI is less likely to have an impact?
- (3) What does an AI-first business mean to you?
- (4) Do you think AI can replace human jobs? Which sectors or job types do you think are most vulnerable?
- (5) What is the difference between AI and ML?

Session 2: Supervised and Unsupervised Algorithms

Learning Material:

- (a) A Concise Explanation of Learning Algorithms with the Mitchell Paradigm.
<https://www.kdnuggets.com/2018/10/mitchell-paradigm-concise-explanation-learning-algorithms.html>

- (b) Introduction to Machine Learning and Problem Framing.
<https://developers.google.com/machine-learning/problem-framing>
- (c) The Building Blocks of AI, by N. Shnoudeh.
<https://nancyshnoudeh.medium.com/the-building-blocks-of-ai-31e5b3ad0857>
- (d) (*) Algorithms – Chapter 6 (Deep Learning), Panos Louridas, MIT Press (2020)
- (e) Recall, Precision, F1, ROC, AUC, and everything. By O. Shalev.
<https://medium.com/swlh/recall-precision-f1-roc-auc-and-everything-542aedf322b9>
- (f) (*) (Optional) Deep Learning – Book (Chapter 1 and 2). <https://mitpress.mit.edu/9780262537551/deep-learning/>
- (g) (Optional, read after class) Evaluating a Machine Learning Model, by J. Jordan.
<https://www.jeremyjordan.me/evaluating-a-machine-learning-model/>

Preparation Questions:

- (1) What are the essential elements of ML models? What purpose does each element serve?
- (2) What are the choices made by designers in developing AI models?
- (3) What is the overfitting and underfitting problem?
- (4) What is the bias variance tradeoff in ML models?
- (5) How should ML models be evaluated and why?
- (6) How should a manager / decision-maker choose from among ML models?
- (7) What are the main differences between deep learning and traditional machine learning models? Can deep neural nets do things that shallow nets cannot? What are these?

Session 3: Generative Models and Reinforcement Learning

There are some technical learning materials below. I'd suggest you skim through these, it's perfectly ok even if you do not fully follow everything in them.

Learning Material:

- (a) Reinforcement Learning Made Simple – Part 1. By Ketan Doshi. Towards Data Science.
<https://towardsdatascience.com/reinforcement-learning-made-simple-part-1-intro-to-basic-concepts-and-terminology-1d2a87aa060>
- (b) (Read Chapter 1, only parts 1.1 to 1.4) Book Chapter from Reinforcement Learning by, R. Sutton and A. Barto, available online at <http://www.incompleteideas.net/sutton/book/first/chapter1.pdf>
- (c) Reinforcement Learning in business: A sneak-peak on the applications of one the most promising AI methods. By Cinelli. (2023)
<https://medium.com/ordina-data/reinforcement-learning-in-business-a-sneak-peak-on-the-applications-of-one-the-most-promising-ai-d9333c77f62d>

- (d) It's time for businesses to chart a course for reinforcement learning. By, J. Corbo, O. Fleming, and N. Hohn, McKinsey & Co.
<https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/its-time-for-businesses-to-chart-a-course-for-reinforcement-learning>
- (e) (*) Algorithms – Chapter 6 (Deep Learning), Panos Louridas, MIT Press (2020)
- (f) What is a Generative Model?
<https://developers.google.com/machine-learning/gan/generative>
- (g) (Technical, Optional) Large Language Models (1-hour talk). A. Karpathy.
https://www.youtube.com/watch?v=zjkBMFhNj_g
- (h) (Technical, Optional) Understanding Variational Autoencoders (VAEs). J. Rocca.
<https://towardsdatascience.com/understanding-variational-autoencoders-vaes-f70510919f73?gi=62a0e0bd28ee>
- (i) (Optional, Read after class) A generative AI reset: Rewiring to turn potential into value in 2024. By Eric Lamarre, Alex Singla, Alexander Sukharevsky, and Rodney Zimmel. McKinsey Quarterly (2024)
<https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/a-generative-ai-reset-rewiring-to-turn-potential-into-value-in-2024>
- (j) (*) (Optional, Read after class) Generative AI and the Future of Work.
- (k) (*) (Optional, Read Chapter 1 and 5 after class) Generative AI. HBR Collection.
- (l) (*) (Optional, Read after class) Generative AI: What Leaders Need to Know.

Preparation Questions:

- (1) What are the use cases for reinforcement learning?
- (2) Can the same problem be solved using both supervised and reinforcement learning? Between unsupervised and reinforcement learning? If so, how would you decide which approach to use?
- (3) How does a generative model differ from a discriminative model?
- (4) Is a generative model limited to exploring only the scope of the training data, or can it generate examples outside of the training data?
- (5) What is common to large language models (LLMs) and VAEs? What are some differences between these models?
- (6) (How) Do you think Generative AI will impact your career over the next few years?
- (7) Do you think there are uniquely human things that Generative AI would *never* be able to do as well as humans?

Session 4: Economics of AI

Learning Material:

- (a) (*) Prediction Machines Book – Chapters 3 and 7.
https://www.google.com/books/edition/Prediction_Machines/wJY4DwAAQBAJ
- (b) (*) The (Surprisingly) Simple Economics of Artificial Intelligence, by A. Agrawal, J. Gans and A. Goldfarb, HBR.
- (c) Reimagining Your Business for AI, by T. Fountaine, B. McCarthy, and T. Saleh, McKinsey.
<https://www.mckinsey.com/business-functions/mckinsey-analytics/our-insights/reimagining-your-business-for-ai>

Preparation Questions:

- (1) Will machines disrupt humans in prediction questions? Why or why not?
- (2) How do you convert a non-prediction problem to a prediction problem?
- (3) Would achieving a high prediction accuracy transform your business? How?
- (4) What are some complements to prediction that we should focus on?

Session 5: The AI Human Interface

Read a, b, f and h before class and the others after class.

Learning Material:

- (a) (*) What AI Driven Decision Making Looks Like, by E. Colson, HBR.
- (b) (*) AI Should Augment Human Intelligence Not Replace It, by D. De Cremer and G. Kasparov, HBR
- (c) The Human Factor in AI Based Decision Making, by P. Meissner and C. Keding, MIT Sloan Review.
<https://sloanreview.mit.edu/article/the-human-factor-in-ai-based-decision-making/>
- (d) (*) How to Win with Machine Learning, by A. Agrawal, J. Gans, and A. Goldfarb, HBR
- (e) (*) When Machine Learning Goes off the Rails, by B. Babic, I. Cohen, T. Evgeniou, & S. Gerke, HBR <https://hbsp.harvard.edu/tu/cfcb7051>
- (f) Explaining Explanations: An Overview of Interpretability of Machine Learning, by L. H. Gilpin, D. Bau, B. Z. Yuan, A. Bajwa, M. Specter and L. Kagal, Arxiv.
<https://arxiv.org/pdf/1806.00069.pdf>
- (g) (Book) Interpretable Machine Learning: A Guide for Making Black Box Models Explainable by Christoph Molnar, Chapter 3 and 5.
<https://christophm.github.io/interpretable-ml-book/>

Preparation Questions:

- (1) Will machines disrupt humans in prediction questions? Why or why not?
 - (2) What is the role of interpretability in ML? What happens when you don't have interpretability?
 - (3) Is explainability different from interpretability? Explain
 - (4) How and when should domain knowledge be incorporated into ML algorithms or the AI decision making process more broadly?
 - (5) What are the risks of AI?
-

Session 6: Ethical Issues in AI

Learning Material:

- (a) (*) The Ethical Algorithm (Book) – Chapter 2
- (b) (*) Introduction and Chapter 2. The Alignment Problem: Machine Learning and Human Values. B. Christian. New York: WW Norton & Company.
- (c) How Machine Learning Pushes Us to Define Fairness
<https://hbr.org/2019/11/how-machine-learning-pushes-us-to-define-fairness>
- (d) What Are Deepfakes and How Are They Created?. IEEE Spectrum (2020)
<https://spectrum.ieee.org/what-is-deepfake>
- (e) (Optional) Fair ML Book (read Chapter 2: Introduction)
<https://fairmlbook.org/pdf/fairmlbook.pdf>
- (f) (Optional) News articles:
 - New Era of AI Deepfakes Complicates 2024 Elections. Robert McMillan, Alexa Corse and Dustin Volz. WSJ (2024).
<https://www.wsj.com/tech/ai/new-era-of-ai-deepfakes-complicates-2024-elections-aa529b9e>
 - OpenAI's GPT Is a Recruiter's Dream Tool. Tests Show There's Racial Bias. By Leon Yin, Davey Alba and Leonardo Nicoletti. Bloomberg.
<https://www.bloomberg.com/graphics/2024-openai-gpt-hiring-racial-discrimination/>
 - Amazon scraps secret AI recruiting tool that showed bias against women. By J. Dastin, Reuters.
<https://www.reuters.com/article/us-amazon-com-jobs-automation-insight/amazon-scraps-secret-ai-recruiting-tool-that-showed-bias-against-women-idUSKCN1MK08G>
 - Machine Bias and Risk Assessments in Criminal Sentencing, by J. Angwin, J. Larson, S. Mattu and L. Kirchner, ProPublica.
<https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

- (Audio) How a computer scientist fights bias in algorithms. By V. Stollberger, NPR.
<https://www.npr.org/2022/03/14/1085160422/computer-science-inequality-bias-algorithms-technology>
- Inside the AI Factory: the humans that make tech seem human. by Josh Dzieza. The Verge (2023) <https://www.theverge.com/features/23764584/ai-artificial-intelligence-data-notation-labor-scale-surge-remotasks-openai-chatbots>
- The Deeper Problem With Google's Racially Diverse Nazis. Chris Gilliard. The Atlantic (2024)
<https://www.theatlantic.com/technology/archive/2024/02/google-gemini-diverse-nazis/677575/>

Preparation Questions:

- (1) In the articles about Amazon gender bias and racial bias in criminal sentencing, what were the underlying reasons for bias?
- (2) What are the different notions of fairness that you can consider? Which ones do you think we should use and why?
- (3) Are these ideas of fairness consistent with one another or potentially in conflict with one another?
- (4) What are the pros and cons of considering protected group membership in making a decision? Are there situations in which such consideration is legal, and situations in which it is illegal?
- (5) Consider each of the different steps in the ML pipeline or process that we have examined. For each step, identify whether bias can be introduced or not. Explain your reasoning.