Research and Teaching Overview

Vineet Kumar

Yale School of Management

Presenting at:
Virginia Tech – Pamplin School of Business
August 2025

Research Overview

Digital Economy – Business Models

Role of Human Knowledge in Research

ML has typically been atheoretical

- One view of ML – advanced form of *statistical pattern matching*
  - Similar model (CNN) used both for detecting lung cancer (medicine) and for detecting stars (astronomy)

My Take

Our **domain knowledge (theory)** has a lot to add

Role of Human Knowledge in Research

ML has typically been atheoretical

Why add domain knowledge?

Can improve predictive *accuracy*, *explainability*, provide *guarantees*

Role of Human Knowledge in Research

ML has typically been atheoretical

Why add domain knowledge?

Can improve predictive *accuracy*, *explainability*, provide *guarantees*

- A Theory-Based Explainable Deep Learning Architecture for Music Emotion. *Marketing Science* 44 (1), 196-219

Role of Human Knowledge in Research

ML has typically been atheoretical

Why add domain knowledge?

Can improve predictive *accuracy*, *explainability*, provide *guarantees*

- A Theory-Based Explainable Deep Learning Architecture for Music Emotion. *Marketing Science* 44 (1), 196-219
- Generative Interpretable Visual Design: Using Disentanglement for Visual Conjoint Analysis. *Journal of Marketing Research* 62.3 (2025): 405-428

Role of Human Knowledge in Research

ML has typically been atheoretical

Why add domain knowledge?

Can improve predictive *accuracy*, *explainability*, provide *guarantees*

- A Theory-Based Explainable Deep Learning Architecture for Music Emotion. *Marketing Science* 44 (1), 196-219
- Generative Interpretable Visual Design: Using Disentanglement for Visual Conjoint Analysis. *Journal of Marketing Research* 62.3 (2025): 405-428
- Nonparametric Bandits Leveraging Informational Externalities to Learn the Demand Curve. Forthcoming at *Marketing Science*

Role of Human Knowledge in Research

ML has typically been atheoretical

Why add domain knowledge?

Can improve predictive *accuracy*, *explainability*, provide *guarantees*

- A Theory-Based Explainable Deep Learning Architecture for Music Emotion. *Marketing Science* 44 (1), 196-219
- Generative Interpretable Visual Design: Using Disentanglement for Visual Conjoint Analysis. *Journal of Marketing Research* 62.3 (2025): 405-428
- Nonparametric Bandits Leveraging Informational Externalities to Learn the Demand Curve. Forthcoming at *Marketing Science*
- Market Structure Mapping with Visual Characteristics. (Research in progress)

Generative Interpretable Visual Design

Sisodia, Burnap and Kumar

Presenting at:
Virginia Tech – Pamplin School of Business
August 2025

Visual (or aesthetic) design matters across many product categories . . .

Visual (or aesthetic) design matters across many product categories . . .

Cars

Fashion

Visual (or aesthetic) design matters across many product categories . . .

Cars

Fashion

Furniture

. . . even for mundane categories like yogurt

"We worked hard to get the packaging right . . . American yogurt has always been sold in containers with relatively narrow openings. In Europe yogurt containers are wider and squatter, and that's what I wanted for Chobani."

—Hamdi Ulukaya, Founder & CEO, Chobani

Consumer Preferences for Visual Design

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

What about preferences in visual space?

Cannot do this because characteristics for visual design are unknown!

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

What about preferences in visual space?

Demand Estimation: Big Picture

Goal:

Obtain consumer preferences for visual design (conjoint or market data)

Demand Estimation for Products in Differentiated Product Markets in Economics and Marketing

- Builds on foundation of Lancaster (1966), Kotler (1967)
- Products are bundles or collections of **characteristics**
- Preferences over products $\implies$ Preferences over characteristics makes this problem feasible.
- Car is represented as: (mpg, horsepower, range) etc.

What about preferences in visual space?

Cannot do this because characteristics for visual design are unknown!

What this research seeks to do

Research Goals

Obtain **human-interpretable** visual characteristics (not outliers) directly from unstructured product image data:

- *automatically discover* and extract characteristics for products

What this research seeks to do

Research Goals

Obtain **human-interpretable** visual characteristics (not outliers) directly from unstructured product image data:

- *automatically discover* and extract characteristics for products
- *quantify* these characteristics

What this research seeks to do

Research Goals

Obtain **human-interpretable** visual characteristics (not outliers) directly from unstructured product image data:

- *automatically discover* and extract characteristics for products
- *quantify* these characteristics
- *generate* visual designs that span the space of visual characteristics

What this research seeks to do

Research Goals

Obtain **human-interpretable** visual characteristics (not outliers) directly from unstructured product image data:

- *automatically discover* and extract characteristics for products
- *quantify* these characteristics
- *generate* visual designs that span the space of visual characteristics

What this research seeks to do

Research Goals

Obtain **human-interpretable** visual characteristics (not outliers) directly from unstructured product image data:

- *automatically discover* and extract characteristics for products
- *quantify* these characteristics
- *generate* visual designs that span the space of visual characteristics

Hyundai: $(3, 8, 5, 9)$ compared to BMW: $(1, 3, 10, 1)$

Research Goals

Hyundai: $(3, 8, 5, 9)$ compared to BMW: $(1, 3, 10, 1)$
Several questions come to mind:

- What does the first number represent? Does 3 mean something different from 1?
- Can humans interpret these numeric values?
- What domain knowledge does the model need to have?

Why Visual Characteristics?

Modeling Visual Characteristics: A comparison of methods

Modeling Visual Characteristics: A comparison of methods

Modeling Visual Characteristics: A comparison of methods

Automatic + Uninterpretable

- A - Bajari, P. L. et al. (2021) : Hedonic prices and quality adjusted price indices powered by AI, *CENMAP working paper*
- B - Law, S., et al. (2019) : Take a look around: using street view and satellite images to estimate house prices. *ACM Transactions on Intelligent Systems and Technology (TIST)*
- C - Aubry, S., et al. (2019) : Machine learning, human experts, and the valuation of real assets. *CFS Working Paper Series*

Modeling Visual Characteristics: A comparison of methods

Automatic + Uninterpretable

- A - Bajari, P. L. et al. (2021) : Hedonic prices and quality adjusted price indices powered by AI, *CENMAP working paper*
- B - Law, S., et al. (2019) : Take a look around: using street view and satellite images to estimate house prices. *ACM Transactions on Intelligent Systems and Technology (TIST)*
- C - Aubry, S., et al. (2019) : Machine learning, human experts, and the valuation of real assets. *CFS Working Paper Series*

Manual + Interpretable

- 1 - Zhang, M. et al. (2022) : Can consumer-posted photos serve as a leading indicator of restaurant survival? Evidence from yelp. *Management Science*
- 2 - Liu, Y., et al. (2017) : The effects of products' aesthetic design on demand and marketing-mix effectiveness: The role of segment prototypicality and brand consistency. *Journal of Marketing*
- 3 - Zhang, S., et al. (2021) : What makes a good image? Airbnb demand analytics leveraging interpretable image features. *Management Science*

What is disentanglement?

Bengio et al (2013)

"A disentangled representation can be defined as one where **single latent units** *are sensitive to changes in single generative factors, while being relatively invariant to changes in other factors"*

What is disentanglement?

Bengio et al (2013)

> *"A disentangled representation can be defined as one where* **single latent units** *are sensitive to changes in single generative factors, while being relatively invariant to changes in other factors"*

- Latent Units ($\mathbf{z}$): Dimensions in the model's latent space
- Generative factors ($\mathbf{c}$): Human-interpretable true characteristics

Idea: Reality or Data generating process is compositional based on generative factors.

What is disentanglement?

Stick

What is disentanglement?

Bengio et al (2013)

"A disentangled representation can be defined as one where **single latent units** *are sensitive to changes in single generative factors, while being relatively invariant to changes in other factors"*

What is disentanglement?

Bengio et al (2013)

"A disentangled representation can be defined as one where **single latent units** *are sensitive to changes in single generative factors, while being relatively invariant to changes in other factors"*

- Latent Units ($\mathbf{z}$): What algorithm discovers – dimensions in the model's latent space
- Generative factors ($\mathbf{c}$): Human-interpretable

What is disentanglement?

Bengio et al (2013)

"A disentangled representation can be defined as one where **single latent units** *are sensitive to changes in single generative factors, while being relatively invariant to changes in other factors"*

- Latent Units ($\mathbf{z}$): What algorithm discovers – dimensions in the model's latent space
- Generative factors ($\mathbf{c}$): Human-interpretable

Goal: One to one mapping between $z \Leftrightarrow c$

Product Images and Parts of Watch

Disentangled and Entangled Representations

Example of *Entangled* Visual Characteristics

Entangled
Dial Color and
Strap Color

Example of *Disentangled* Visual Characteristics

Dial Color

Strap Color

Models in Existing Literature

Model	Goal
Autoencoder (AE)	Reconstruction accuracy
Variational Autoencoder (VAE)	...+ structured latent space
Disentanglement	...+ ...+ statistically independent latent space

Roadmap of Our Research Approach

Contribution

We aim to overcome this impossibility theorem with a simple approach of using structured product characteristics.

Impossibility Theorem

Impossibility Theorem

Unsupervised (*i.e. only images*) learning of disentangled representations is *fundamentally impossible* except under certain restrictive conditions.[a]

[a] Locatello, Francesco, et al. "Challenging common assumptions in the unsupervised learning of disentangled representations." ICML. PMLR, 2019.

Implication: Every disentangled representation can have other *infinite* equivalent entangled representations.

ML Approach to Impossibility Theorem

Impossibility: Without Supervision, every disentangled representation can have other *infinite* equivalent entangled representations.

- ML researchers recognize the challenge of impossibility

ML Approach to Impossibility Theorem

Impossibility: Without Supervision, every disentangled representation can have other *infinite* equivalent entangled representations.

- ML researchers recognize the challenge of impossibility
- Need a supervisory signal

ML Approach to Impossibility Theorem

Impossibility: Without Supervision, every disentangled representation can have other *infinite* equivalent entangled representations.

- ML researchers recognize the challenge of impossibility
- Need a supervisory signal
- ML methods assume that ground truth is known by researchers

ML Approach to Impossibility Theorem

Impossibility: Without Supervision, every disentangled representation can have other *infinite* equivalent entangled representations.

- ML researchers recognize the challenge of impossibility
- Need a supervisory signal
- ML methods assume that ground truth is known by researchers
  - Human labeling

ML Approach to Impossibility Theorem

Impossibility: Without Supervision, every disentangled representation can have other *infinite* equivalent entangled representations.

- ML researchers recognize the challenge of impossibility
- Need a supervisory signal
- ML methods assume that ground truth is known by researchers
  - Human labeling
- Can we use this approach to discover visual characteristics?

Impossibility Theorem – Implications

Impossibility Theorem – Implications

Common approach to ground truth in ML is to get humans to label[1]

What's the Problem?

- Ground truth on visual characteristics is *unknown*.

[1] Locatello, Francesco, et al. "Disentangling factors of variation using few labels." ICLR. 2020.

Impossibility Theorem – Implications

Common approach to ground truth in ML is to get humans to label[1]

What's the Problem?

- Ground truth on visual characteristics is *unknown*.
- Need to ensure humans understand what these labels are and *how to quantify them* for each image

[1] Locatello, Francesco, et al. "Disentangling factors of variation using few labels." ICLR. 2020.

Schematic of Proposed Approach

Contribution

- **Solution** without ground truth on visual characteristics:
- Leverage **structured product characteristics** to provide a supervisory signal for disentanglement

Model

- Learn model parameters by minimizing loss $L(\theta, \phi; \mathbf{x}, \mathbf{z})$ of integrated model
- θ and ϕ are encoder and decoder parameters; $\mathbf{x}$ are images

$$\underbrace{L(\theta, \phi, \mathbf{w}); \mathbf{x}, \mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z}, \mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z}) || \prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+ \gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j) || p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}), \mathbf{y})}_{\text{Supervised Loss}}$$

Model

- Learn model parameters by minimizing loss $L(\theta, \phi; \mathbf{x}, \mathbf{z})$ of integrated model
- θ and ϕ are encoder and decoder parameters; $\mathbf{x}$ are images

$$\underbrace{L(\theta, \phi, \mathbf{w}); \mathbf{x}, \mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z}, \mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z}) || \prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+ \gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j) || p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}), \mathbf{y})}_{\text{Supervised Loss}}$$

Loss Term	Why is this term included?
Reconstruction	Promotes accurate reconstruction of images
Mutual Information	Minimizes redundant information
Total Correlation	Promotes statistical independence between visual characteristics
Dimension-Wise KL	Penalizes deviations from a prior
Supervised	Provides a signal to address the impossibility theorem

Model – Role of Supervised Loss

$$\underbrace{L(\theta,\phi,\mathbf{w});\mathbf{x},\mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z},\mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z})||\prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+\gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j)||p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}),\mathbf{y})}_{\text{Supervised Loss}}$$

- Supervised Loss is used to predict signal from latent representation z: $s = f(z)$

Model – Role of Supervised Loss

$$\underbrace{L(\theta,\phi,\mathbf{w});\mathbf{x},\mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z},\mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z})||\prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+\gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j)||p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}),\mathbf{y})}_{\text{Supervised Loss}}$$

- Supervised Loss is used to predict signal from latent representation z: $s = f(z)$
- Can use structured product characteristics as signals: brand, price, material etc.

Model – Role of Supervised Loss

$$\underbrace{L(\theta,\phi,\mathbf{w});\mathbf{x},\mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z},\mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z})||\prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+\gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j)||p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}),\mathbf{y})}_{\text{Supervised Loss}}$$

- Supervised Loss is used to predict signal from latent representation z: $s = f(z)$
- Can use structured product characteristics as signals: brand, price, material etc.

Model – Role of Supervised Loss

$$\underbrace{L(\theta,\phi,\mathbf{w});\mathbf{x},\mathbf{z})}_{\text{Total Loss}} = \underbrace{\mathbf{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}|\mathbf{z})\right]}_{\text{Reconstruction Loss}} + \alpha \underbrace{I_q(\mathbf{z},\mathbf{x})}_{\text{Mutual Information Loss}} + \beta \underbrace{KL\left[q(\mathbf{z})||\prod_{j=1}^{J} q(z_j)\right]}_{\text{Total Correlation Loss}}$$

$$+\gamma \underbrace{\sum_{j=1}^{J} KL\left[q(z_j)||p(z_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \delta \underbrace{P(\hat{\mathbf{y}}(\mathbf{z}),\mathbf{y})}_{\text{Supervised Loss}}$$

- Supervised Loss is used to predict signal from latent representation z: $s = f(z)$
- Can use structured product characteristics as signals: brand, price, material etc.

Idea to Overcome Impossibility Theorem

If the supervisory signal is sufficiently correlated with visual characteristics, then it can help obtain the unique (true) disentangled representation

Why might brand aid the disentanglement model?

Brand as a Supervisory Signal

Idea: Brands have a specific "look" that can be correlated with visual appearance (and therefore visual characteristics)

Evaluating Visual Characteristics

Visual Characteristics: Human Interpretable?

- Are these visual characteristics human interpretable?

Visual Characteristics: Human Interpretable?

- Are these visual characteristics human interpretable?

Starting from the image on the left, **what part of the watch changes the most** as you go from left to right? Carefully check both large and small visual aspects. Go through each part of the watch one by one before selecting any option. Refer to the above image to see parts of the watch.

Note: Images are low-quality on purpose

- ○ Bezel
- ○ Crown
- ○ Date Window
- ○ Dial
- ○ Hands
- ○ Hour Marker
- ○ Lug
- ○ Strap

How is that part of the watch changing?

Visual Characteristics: Quantification?

Interpretability and Quantification

Visual characteristic	Interpretability Survey	Quantification Survey
Dial Size	76%	83%
Dial Color	80%	92%
Strap Color	88%	92%
Rim (Bezel) Color	79%	88%
Dial Shape	87%	68%
Knob (Crown) Size	70%	85%

Discovered Visual characteristics

Discovered Visual characteristics

Discovered Visual characteristics

Discovered Visual characteristics

Discovered Visual characteristics

Discovered Visual characteristics

Density of Discovered Visual characteristics

Density of Discovered Visual characteristics

Research Goals

- Visual conjoint has been challenging to do. With disentanglement we can create counterfactual designs to span the space.

Conjoint Model Accuracy

Generated Watches

Model	Out-of-Sample Hit Rate (SD)
Disentangled Embedding + Logit Model (–)	63.16% (2.34%)
Disentangled Embedding + Neural Net (–)	65.81% (2.22%)
Pretrained Deep Learning Model Embedding (O)	68.31% (1.54%)
Disentangled Embedding + Neural Net (O)	67.52% (0.92%)
Disentangled Embedding + Random Forest (O)	68.77% (0.90%)
Disentangled Embedding + XGBoost (O)	69.10% (0.41%)
Disentangled Embedding + HB Model (O + U)	**71.61% (1.87%)**
Disentangled Embedding + HB Model + Interactions (O + U)	70.68% (1.35%)

- Pretrained Deep learning model is trained on *millions of images*, and has millions of parameters
- Our Hierarchical Bayes (HB) model has a small number parameters, and all predictions are based on only 6 visual characteristics
- With 6 visual characteristics, our HB model outperforms the pretrained deep neural net

Conclusion

We obtain interpretable visual characteristics directly from unstructured product images

- *automatically discover* (extract) characteristics

Conclusion

We obtain interpretable visual characteristics directly from unstructured product images

- *automatically discover* (extract) characteristics
- *quantify* these characteristics

Conclusion

We obtain interpretable visual characteristics directly from unstructured product images

- *automatically discover* (extract) characteristics
- *quantify* these characteristics
- *generate* visual design that span the space of visual characteristics

Conclusion

We obtain interpretable visual characteristics directly from unstructured product images

- *automatically discover* (extract) characteristics
- *quantify* these characteristics
- *generate* visual design that span the space of visual characteristics

Conclusion

We obtain interpretable visual characteristics directly from unstructured product images

- *automatically discover* (extract) characteristics
- *quantify* these characteristics
- *generate* visual design that span the space of visual characteristics

Applications

We then used the model to:

- generate new counterfactual designs to obtain consumer preferences over visual characteristics.
- obtain ideal point visual designs corresponding to different consumer segments

Market Structure Mapping with Interpretable Visual Characteristics

Sisodia and Kumar

Presenting at:
Virginia Tech – Pamplin School of Business
August 2025

What is a Market Structure Map?

A common element of market structure analysis is the derivation of a **market structure map**, *that is, a spatial representation of firms' competitive positions relative to one another based on some measure of their competitive relationships (DeSarbo et al. 1993)*

- Graphical representations of product positioning (e.g. Multidimensional Scaling)

What is a Market Structure Map?

A common element of market structure analysis is the derivation of a **market structure map**, *that is, a spatial representation of firms' competitive positions relative to one another based on some measure of their competitive relationships (DeSarbo et al. 1993)*

- Graphical representations of product positioning (e.g. Multidimensional Scaling)
- Strategic tool for understanding competitive landscape that enable us to:

What is a Market Structure Map?

A common element of market structure analysis is the derivation of a **market structure map**, *that is, a spatial representation of firms' competitive positions relative to one another based on some measure of their competitive relationships (DeSarbo et al. 1993)*

- Graphical representations of product positioning (e.g. Multidimensional Scaling)
- Strategic tool for understanding competitive landscape that enable us to:
  - Identify market gaps

What is a Market Structure Map?

A common element of market structure analysis is the derivation of a **market structure map**, *that is, a spatial representation of firms' competitive positions relative to one another based on some measure of their competitive relationships (DeSarbo et al. 1993)*

- Graphical representations of product positioning (e.g. Multidimensional Scaling)
- Strategic tool for understanding competitive landscape that enable us to:
  - Identify market gaps
  - Spot new opportunities

What is a Market Structure Map?

A common element of market structure analysis is the derivation of a **market structure map**, *that is, a spatial representation of firms' competitive positions relative to one another based on some measure of their competitive relationships (DeSarbo et al. 1993)*

- Graphical representations of product positioning (e.g. Multidimensional Scaling)
- Strategic tool for understanding competitive landscape that enable us to:
  - Identify market gaps
  - Spot new opportunities
  - Identify competitors and evaluate positioning

Market Structure Map: Examples

Compositional (Left) & Decompositional (Right)

Lee, Sudhir and Steckel (2002), *Journal of Marketing Research*

Erdem (1996), *Marketing Science*

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices
- Sales data (cross-elasticity of demand)

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices
- Sales data (cross-elasticity of demand)
- User comments and reviews

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices
- Sales data (cross-elasticity of demand)
- User comments and reviews
- Search Co-occurrence

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices
- Sales data (cross-elasticity of demand)
- User comments and reviews
- Search Co-occurrence

Market Structure Mapping

Sources of Data

Compositional (Building up)

- Product characteristics – functional, psychological

Decompositional (Breaking down)

- Survey data by asking 1st and 2nd choices
- Sales data (cross-elasticity of demand)
- User comments and reviews
- Search Co-occurrence

What's common?

Neither of these approaches have considered visual characteristics.

Are these cars close substitutes?

First, consider *functional* product characteristics:

(A)	**(B)**
\$179,000-\$243,000, 553 HP, 16-23 MPG	From \$242,700, 542-650 HP, 14-22 MPG

Are these cars close substitutes?

First, consider *functional* product characteristics:

(A)

$179,000-$243,000, 553 HP, 16-23 MPG

Ferrari California

(B)

From $242,700, 542-650 HP, 14-22 MPG

Bentley Continental GT

Importance of Visual Design

Functional characteristics alone are *not sufficient* to even place the car in an appropriate segment (submarket)

Visual design matters *a lot* for cars

"Exterior look/design is the top reason shoppers avoid a particular vehicle (30%), followed by cost (17%)."

–JD Power Avoider Study 2015

Challenges in using visual images for market structure mapping

Three Challenges:

1. Inputs from Raw Images: Backgrounds, lighting, and camera angles overshadow actual product design, causing mis-clusters.

Challenges in using visual images for market structure mapping

Three Challenges:

1. Inputs from Raw Images: Backgrounds, lighting, and camera angles overshadow actual product design, causing mis-clusters.
2. Scalability: Need to work with hundreds or thousands of product models / images, so challenging to have human intervention

Challenges in using visual images for market structure mapping

Three Challenges:

1. Inputs from Raw Images: Backgrounds, lighting, and camera angles overshadow actual product design, causing mis-clusters.
2. Scalability: Need to work with hundreds or thousands of product models / images, so challenging to have human intervention
3. Black-Box (Un)interpretable Embeddings: Pretrained networks may focus on brand logos or color, complicating interpretation of true style.

Challenges in using visual images for market structure mapping

Three Challenges:

1. Inputs from Raw Images: Backgrounds, lighting, and camera angles overshadow actual product design, causing mis-clusters.
2. Scalability: Need to work with hundreds or thousands of product models / images, so challenging to have human intervention
3. Black-Box (Un)interpretable Embeddings: Pretrained networks may focus on brand logos or color, complicating interpretation of true style.
   - We need to be able to choose what visual characteristics matter for market structure

Challenges in using visual images for market structure mapping

Three Challenges:

1. Inputs from Raw Images: Backgrounds, lighting, and camera angles overshadow actual product design, causing mis-clusters.
2. Scalability: Need to work with hundreds or thousands of product models / images, so challenging to have human intervention
3. Black-Box (Un)interpretable Embeddings: Pretrained networks may focus on brand logos or color, complicating interpretation of true style.
   - We need to be able to choose what visual characteristics matter for market structure
   - Does color matter? Does angle matter?

Should Color be included in Visual Characteristics Maps?

- Should a red Ferrari California be close substitute to a red Toyota Corolla?

Data

- **UK Automobile Market**
  - 2,439 make-model-year observations from 2008 to 2017
  - 379 unique models from 49 unique makes
- **Sales and Auto Characteristics:**
  - Characteristics (price, MPG, horsepower, weight, dimensions, etc.) from Parker's
  - Sales from DVM-CAR[2]
- **Images:**
  - Front-facing automobile images from DVM-CAR
  - Converted from color to grayscale to *focus on shape* rather than color.
  - (very) Low-resolution: 128×128 pixels

[2] Huang, Jingmin, et al. (2022), *IEEE International Conference on Big Data*

Segments in the UK automobile market

Segment	Description
A	Minicars
B	Subcompact
C	Compact
D	Mid-size
E	Mid-size Luxury
J	SUV
M	MPV

Seg A (Minicar)

Seg B (Subcompact)

Seg C (Compact)

Seg D (Mid-Sized)

Seg E (Mid-Sized Luxury)

Seg J (SUV)

Seg M (MPV)

Summary Statistics of the 2013 UK Auto Market

Variable	Mean	St. Dev.	Min.	Max.
Total Quantity Sold	8,074.834	13,714.100	1	113,390
Price (in £000 units)	26.333	14.668	7.868	108.624
MPG (tens of miles per gallon)	4.999	1.058	2.250	7.200
Weight (in 10 lbs)	327.704	2.444	324.506	332.106
HP/Wt (in HP per 10 lbs.)	0.461	0.169	0.060	1.347
Length (in 1000 inches)	1.724	0.178	1.062	2.054
Width (in 1000 inches)	0.756	0.063	0.580	0.899
Height (in 1000 inches)	0.616	0.054	0.537	0.780
Wheelbase (in 1000 inches)	1.046	0.081	0.735	1.266
Space (length $\times$ width)	1.310	0.215	0.697	1.759
Vehicle Segment (Proportion):				
Segment A (Minicars)	0.112	0.316	0	1
Segment B (Subcompact)	0.120	0.326	0	1
Segment C (Compact)	0.162	0.369	0	1
Segment D (Mid-Size)	0.129	0.335	0	1
Segment E (Mid-Size Luxury)	0.071	0.257	0	1
Segment J (SUV)	0.241	0.428	0	1
Segment M (MPV)	0.166	0.373	0	1
Country of Origin (Proportion):				
France	0.108	0.311	0	1
Germany	0.241	0.428	0	1
Japan	0.220	0.415	0	1
South Korea	0.091	0.289	0	1
UK	0.108	0.311	0	1
USA	0.058	0.234	0	1

Match each model with front-facing image

Is this sufficient?

"Market research studies have shown that 70% of consumers identify and judge automobiles by the appearance of headlights and grille located on the face of the automobile."

– The Wall Street Journal 2006

Loss Function & Supervisory Signals

$$\underbrace{L(\theta, \phi; \mathbf{m}, \mathbf{v})}_{\text{Disentanglement Loss}} = \underbrace{-\mathbf{E}_{q_\phi(\mathbf{v}|\mathbf{m})}\left[\log p_\theta(\mathbf{m}|\mathbf{v})\right]}_{\text{Reconstruction Loss}} + \underbrace{I_q(\mathbf{v}, \mathbf{m})}_{\text{Mutual Information Loss}} \tag{1}$$

$$+ \lambda_1 \underbrace{KL\left[q(\mathbf{v})|| \prod_{j=1}^{J} q(v_j)\right]}_{\text{Total Correlation Loss}} + \underbrace{\sum_{j=1}^{J} KL\left[q(v_j)||p(v_j)\right]}_{\text{Dimension-Wise KL Divergence Loss}} + \lambda_2 \underbrace{P(\widehat{\mathbf{y}(\mathbf{v})}, \mathbf{y})}_{\text{Supervised Loss}}$$

Number of Signals	Supervisory Signals	λ_1	λ_2	UDR
3	Wheelbase, Width, Height	50	10	0.739
3	HP/Weight, MPG, Space	50	30	0.710
1	Price	50	30	0.708
1	Weight	50	40	0.708
1	Wheelbase	50	30	0.690
1	Width	50	5	0.689
3	Length, Width, Height	50	40	0.678
1	Length	50	40	0.666
0	Unsupervised β-TCVAE	50	0	0.658
1	Height	30	20	0.378
1	Country of Origin	10	10	0.139
1	Segment	10	10	0.134
1	Unsupervised VAE	1	0	0.073
1	Unsupervised AE	0	0	0.074
1	Make	1	1	0.072

$\lambda_1 \in [1, 5, 10, 20, 30, 40, 50]$ and $\lambda_2 \in [0, 1, 5, 10, 20, 30, 40, 50]$.

Visual Characteristic 1: Body Shape

Correlation Between Discovered Visual Dimensions and Physical Vehicle Measures

	Wheelbase	Weight	Length	Height	Width	Height/Width Ratio
Body Shape	0.30	0.33	0.39	-0.28	0.25	-0.42

- Hatchback-like profiles $\rightarrow$ smaller wheelbase, lower length, and lighter
- Hatchback-like profiles $\rightarrow$ taller and narrower

Visual Characteristic 2: Boxiness

Correlation Between Discovered Visual Dimensions and Physical Vehicle Measures

	Wheelbase	Weight	Length	Height	Width	Height/Width Ratio
Boxiness	0.05	-0.07	0.14	-0.59	0.02	-0.49

- Higher degree of boxiness → Taller and upfront from the front
- Lower degree of boxiness → Flatter and sleeker
- Boxiness captures cabin "uprightness" (not related to length or wheelbase)

Visual Characteristic 3: Grille Height

Correlation Between Discovered Visual Dimensions and Physical Vehicle Measures

	Wheelbase	Weight	Length	Height	Width	Height/Width Ratio
Grille Height	0.04	0.02	0.05	-0.04	0.03	-0.05

- Reflects stylistic design choices
- Unrelated to size-related factors

Visual Characteristic 4: Grille Width

Correlation Between Discovered Visual Dimensions and Physical Vehicle Measures

	Wheelbase	Weight	Length	Height	Width	Height/Width Ratio
Grille Width	0.12	0.08	0.12	0.03	0.15	-0.09

- Reflects stylistic design choices
- Unrelated to size-related factors

Visual Characteristic 5: Color

- We find the visual characteristic of *Color*.
- Color should not impact market structure because a car is available in many colors.
- We, therefore, exclude it when we create market structure maps.
- This is only possible because of disentanglement.

Disentanglement obtains four interpretable visual characteristics

Left to Right: Vary one visual characteristic, keeping all others fixed

Human Understanding of Visual Characteristics

- Showed respondents sequence of car images

Q1/4: Look at the below image to see the various parts of a car.

Now, carefully examine each car image below from 1 to 5, going from left to right.
Note: Images are low-quality on purpose. Be sure to see all the images 1 to 5.

1 2 3 4 5

LEFT ⟷ RIGHT

How does the car change the <u>most</u> as you go from image 1 to 5? Go through each part of the car one by one before deciding your response. Write it in a few words.

Human Understanding of Visual Characteristics

- Showed respondents sequence of car images
- Asked respondents open-ended questions about visual changes

Q1/4: Look at the below image to see the various parts of a car.

Now, carefully examine each car image below from 1 to 5, going from left to right.
Note: Images are low-quality on purpose. Be sure to see all the images 1 to 5.

How does the car change the <u>most</u> as you go from image 1 to 5? Go through each part of the car one by one before deciding your response. Write it in a few words.

Human Understanding of Visual Characteristics

- Showed respondents sequence of car images
- Asked respondents open-ended questions about visual changes
- Used LLMs (GPT-4, Claude) to identify *common themes* in respondent responses

Q1/4: Look at the below image to see the various parts of a car.

Now, carefully examine each car image below from 1 to 5, going from left to right.
Note: Images are low-quality on purpose. Be sure to see all the images 1 to 5.

1 2 3 4 5

LEFT ⟷ RIGHT

How does the car change the <u>most</u> as you go from image 1 to 5? Go through each part of the car one by one before deciding your response. Write it in a few words.

Human Understanding of Visual Characteristics

- Showed respondents sequence of car images
- Asked respondents open-ended questions about visual changes
- Used LLMs (GPT-4, Claude) to identify *common themes* in respondent responses
- Example: LLM summarizes that respondents are saying that cars becoming lower, flatter, and wider can be thought of as boxiness

Q1/4: Look at the below image to see the various parts of a car.

Now, carefully examine each car image below from 1 to 5, going from left to right.
Note: Images are low-quality on purpose. Be sure to see all the images 1 to 5.

How does the car change the most as you go from image 1 to 5? Go through each part of the car one by one before deciding your response. Write it in a few words.

Quantification Validation Survey Question

Which pair of cars in your judgment are visually more similar? Carefully check both large and small visual aspects. Do not consider any non-visual features like brand or price.

Left Pair

Right Pair

Correlation Between Different Characteristics

	Functional Characteristics				Visual Characteristics			
	Price	MPG	HP/Weight	Space	Boxiness	Body Shape	Grille Height	Grille Width
Price	1.00							
MPG	-0.60	1.00						
HP/Weight	0.74	-0.48	1.00					
Space	0.67	-0.47	0.36	1.00				
Boxiness	0.06	0.04	0.29	0.09	1.00			
Body Shape	0.50	-0.25	0.54	0.36	0.13	1.00		
Grille Height	0.11	0.03	0.12	0.05	0.04	-0.02	1.00	
Grille Width	0.07	-0.05	0.04	0.15	0.01	-0.12	-0.05	1.00

Does Form Follow Function?
Insight 1: Products close in functional space are differentiated visually

- *Operationalization*: Calculate correlation between distances in functional product characteristics and visual product characteristics for each pair of make-models within each segment

Does Form Follow Function?
Insight 1: Products close in functional space are differentiated visually

- *Operationalization*: Calculate correlation between distances in functional product characteristics and visual product characteristics for each pair of make-models within each segment
- *Result*: Low correlation between functional and visual similarity

Does Form Follow Function?
Insight 1: Products close in functional space are differentiated visually

- *Operationalization*: Calculate correlation between distances in functional product characteristics and visual product characteristics for each pair of make-models within each segment
- *Result*: Low correlation between functional and visual similarity
- *Implication*: Visual can create distinctiveness

Does Form Follow Function?
Insight 1: Products close in functional space are differentiated visually

- *Operationalization*: Calculate correlation between distances in functional product characteristics and visual product characteristics for each pair of make-models within each segment
- *Result*: Low correlation between functional and visual similarity
- *Implication*: Visual can create distinctiveness

Does Form Follow Function?
Insight 1: Products close in functional space are differentiated visually

- *Operationalization*: Calculate correlation between distances in functional product characteristics and visual product characteristics for each pair of make-models within each segment
- *Result*: Low correlation between functional and visual similarity
- *Implication*: Visual can create distinctiveness

Segment	Correlation	Std. Error
A (Minicars)	-0.08	0.05
B (Subcompact)	-0.05	0.05
C (Compact)	-0.05	0.04
D (Mid-size)	0.02	0.04
E (Mid-size Luxury)	0.13	0.10
J (SUV)	0.09	0.03
M (MPV)	0.11	0.04

Market Structure Maps

Market Structure Map using only functional characteristics

Segment	Description
B	Subcompact
D	Mid-size
J	SUV

Market Structure Map using only visual characteristics

Segment	Description
B	Subcompact
D	Mid-size
J	SUV

Market Structure Map using both type of characteristics

Segment	Description
B	Subcompact
D	Mid-size
J	SUV

Insight 2: Visual increases differentiation and helps separate segments

Only Functional Characteristics

24.3% segment overlap

Only Visual Characteristics

32.9% segment overlap

Both Characteristics

8.3% segment overlap

Does Disentanglement Provide Further Insights?

Insight 2D: Which visual characteristics increase differentiation?

- Disentanglement identifies interpretable characteristics contributing to differentiation

Insight 2D: Which visual characteristics increase differentiation?

- Disentanglement identifies interpretable characteristics contributing to differentiation
- Seg D & J overlap on most of the functional characteristics but Seg B is different

Insight 2D: Which visual characteristics increase differentiation?

- Disentanglement identifies interpretable characteristics contributing to differentiation
- Seg D & J overlap on most of the functional characteristics but Seg B is different
- Seg D & J are different in boxiness and grille

Insight 3: Different Visual Strategy Across Brands

Segment J – SUVs

Only Functional Characteristics

Only Visual Characteristics

Insight 3: Different Visual Strategy Across Brands

Segment J – SUVs

Only Functional Characteristics

Only Visual Characteristics

Area Share of a Make in Functional Space & Visual Space (Segment J: SUVs)

Make	Models	Area Share (Functional)	Area Share (Visual)	Ratio
Audi	4	17.90%	2.96%	6.05
BMW	4	6.35%	6.48%	0.98
Jeep	4	9.38%	28.42%	0.33

Insight 3D: Different Visual Strategy Across Brands

Disentanglement identifies interpretable characteristics

- Audi is very tight on some visual characteristics, notably, grille height and grille width $\Longrightarrow$ *distinctive look*

Insight 3D: Different Visual Strategy Across Brands

Disentanglement identifies interpretable characteristics

- Audi is very tight on some visual characteristics, notably, grille height and grille width $\Longrightarrow$ *distinctive look*
- BMW is typically a more "boxy" look

Insight 3D: Different Visual Strategy Across Brands

Disentanglement identifies interpretable characteristics

- Audi is very tight on some visual characteristics, notably, grille height and grille width $\Longrightarrow$ *distinctive look*
- BMW is typically a more "boxy" look
- Jeep varies on most visual characteristics

Insight 3: Different Visual Strategy Across Brands

Audi Q3

Audi Q5

Audi Q7

Audi SQ5

BMW X1

BMW X3

BMW X5

BMW X6

Jeep Cherokee

Jeep Compass

Jeep Grand Cherokee

Jeep Wrangler

Insight 4: Product-Level

Only Functional Characteristics

Only Visual Characteristics

- Looking only at functional characteristics, BMW seems to be positioning X5 and X6 too close (cannibalization risk)

Insight 4: Product-Level

Only Functional Characteristics

Only Visual Characteristics

- Looking only at functional characteristics, BMW seems to be positioning X5 and X6 too close (cannibalization risk)
  - Looking at visual, we can see that they are now separated more

Insight 4: Product-Level

Only Functional Characteristics

Only Visual Characteristics

- Looking only at functional characteristics, BMW seems to be positioning X5 and X6 too close (cannibalization risk)
  - Looking at visual, we can see that they are now separated more
- Positioning BMW X3: Should I focus on Cherokee (which is what functional map suggests)?

Insight 4: Product-Level

Only Functional Characteristics

Only Visual Characteristics

- Looking only at functional characteristics, BMW seems to be positioning X5 and X6 too close (cannibalization risk)
  - Looking at visual, we can see that they are now separated more
- Positioning BMW X3: Should I focus on Cherokee (which is what functional map suggests)?
  - Visual map suggests the Grand Cherokee is closer instead

Insight 5: Leaping to Another Category

Only Functional Characteristics

Only Visual Characteristics

Both Characteristics

Insight 5D: Leaping to Another Category

- Compact cars and compact crossover SUVs differ in boxiness
- Compact cars and compact crossover SUVs overlap across most functional characteristics
- Compact crossover SUVs and compact SUVs are similar in visual characteristics
- Compact crossover SUVs and compact SUVs differ in some

Connecting Consumer Search to Market Structure Maps

Consumer Search $\Longleftrightarrow$ Market structure maps?

We use Google Trends to look for consumers searching pairs of models

- Consumer searching for "Honda Civic" and "Toyota Corolla" more likely to be comparing

$$\log_{10}(Y_{AB}+1) = \beta_0 + \beta_1(|X_A - X_B|) + \beta_2(|V_A - V_B|) + \epsilon_{AB}$$

[3] Raw Google Trend values are scaled by each pair's segment sales share

Consumer Search $\Longleftrightarrow$ Market structure maps?

We use Google Trends to look for consumers searching pairs of models

- Consumer searching for "Honda Civic" and "Toyota Corolla" more likely to be comparing
- Models likely in consideration set

$$\log_{10}(Y_{AB}+1) = \beta_0 + \beta_1(|X_A - X_B|) + \beta_2(|V_A - V_B|) + \epsilon_{AB}$$

[3] Raw Google Trend values are scaled by each pair's segment sales share

Consumer Search $\Longleftrightarrow$ Market structure maps?

We use Google Trends to look for consumers searching pairs of models

- Consumer searching for "Honda Civic" and "Toyota Corolla" more likely to be comparing
- Models likely in consideration set

$$\log_{10}(Y_{AB}+1) = \beta_0 + \beta_1(|X_A - X_B|) + \beta_2(|V_A - V_B|) + \epsilon_{AB}$$

- $\log_{10}(Y_{AB}+1)$ is the log-transformed normalized Google Trends score for the pair (model A, model B) belonging to the same segment[3]

[3] Raw Google Trend values are scaled by each pair's segment sales share

Consumer Search $\Longleftrightarrow$ Market structure maps?

We use Google Trends to look for consumers searching pairs of models

- Consumer searching for "Honda Civic" and "Toyota Corolla" more likely to be comparing
- Models likely in consideration set

$$\log_{10}(Y_{AB}+1) = \beta_0 + \beta_1(|X_A - X_B|) + \beta_2(|V_A - V_B|) + \epsilon_{AB}$$

- $\log_{10}(Y_{AB}+1)$ is the log-transformed normalized Google Trends score for the pair (model A, model B) belonging to the same segment[3]
- (X_A, X_B) are standardized functional characteristics

[3] Raw Google Trend values are scaled by each pair's segment sales share

Consumer Search $\Longleftrightarrow$ Market structure maps?

We use Google Trends to look for consumers searching pairs of models

- Consumer searching for "Honda Civic" and "Toyota Corolla" more likely to be comparing
- Models likely in consideration set

$$\log_{10}(Y_{AB}+1) = \beta_0 + \beta_1(|X_A - X_B|) + \beta_2(|V_A - V_B|) + \epsilon_{AB}$$

- $\log_{10}(Y_{AB}+1)$ is the log-transformed normalized Google Trends score for the pair (model A, model B) belonging to the same segment[3]
- $(X_A,\ X_B)$ are standardized functional characteristics
- $(V_A,\ V_B)$ are standardized visual characteristics

[3] Raw Google Trend values are scaled by each pair's segment sales share

Consumer Search $\Longleftrightarrow$ Market structure maps?

	(1)	(2)	(3)	(4)
Constant	2.778***	9.509***	2.974***	8.871***
	(0.205)	(1.425)	(0.291)	(1.479)
Δ HP/Weight	−2.395	−3.761**	−2.886*	−4.021***
	(1.691)	(1.460)	(1.638)	(1.453)
Δ Space	−1.685	1.200	−1.444	1.431*
	(1.068)	(0.851)	(1.031)	(0.853)
Δ MPG	−0.601***	−0.407**	−0.599***	−0.432***
	(0.199)	(0.163)	(0.196)	(0.165)
Δ Price	−0.214***	−0.188***	−0.203***	−0.194***
	(0.037)	(0.031)	(0.036)	(0.032)
Δ Boxiness			−0.263	0.249
			(0.161)	(0.154)
Δ Bodyshape			0.149	−0.361**
			(0.108)	(0.144)
Δ Grille Height			−0.284***	−0.007
			(0.072)	(0.065)
Δ Grille Width			0.324**	−0.001
			(0.138)	(0.145)
Make Fixed Effects	No	Yes	No	Yes
Observations	306	306	306	306
R^2	0.168	0.590	0.247	0.602
Adjusted R^2	0.157	0.559	0.227	0.566
Note:	*p<0.1; **p<0.05; ***p<0.01			

Consumers search for visually similar cars even after accounting for functional similarity

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments
- Different visual strategy across brands

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments
- Different visual strategy across brands
  - Brand consistency (e.g. Audi)

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments
- Different visual strategy across brands
  - Brand consistency (e.g. Audi)
  - Within-brand differentiation (e.g. Jeep) in design

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments
- Different visual strategy across brands
  - Brand consistency (e.g. Audi)
  - Within-brand differentiation (e.g. Jeep) in design
- Product close in functional space often differentiated visually and the other way around

Conclusion and Takeaways

- We quantify human-interpretable visual design characteristics and obtain market structure maps
  - Method allows us to select only *relevant* visual characteristics
- (Interpretable) Visual characteristics can provide an entirely new set of insights at segment (submarket) level, brand level and product level
- Visual design increases differentiation across segments
- Different visual strategy across brands
  - Brand consistency (e.g. Audi)
  - Within-brand differentiation (e.g. Jeep) in design
- Product close in functional space often differentiated visually and the other way around
- Consumers search for visually similar cars even after accounting for functional similarity